Structural boundary repair: two fresh evaluations

Fixed 128-input pool · 48 unique measured circuits · zero fits · policy and selection frozen before labels

Predeclared evidence gates: NOT PASSED — no promotion

General cohort · 32 fixed circuits

Mean absolute error · LUT4 ↓

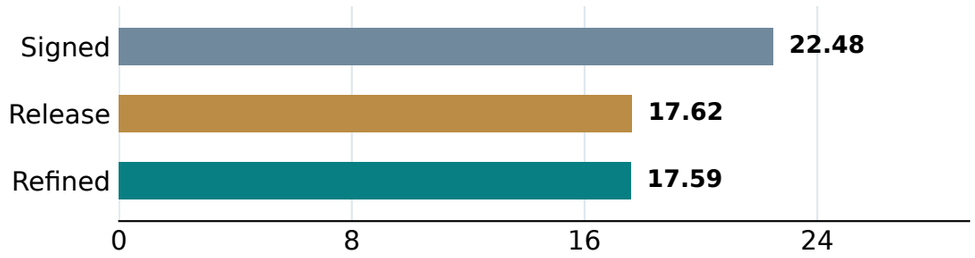

Positive-target MAPE · % ↓

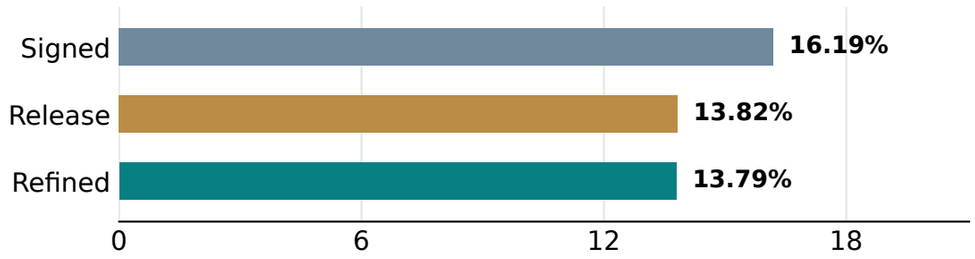

False accepts / false rejects · separate gates

Cap · LUT4	Signed	Release	Refined
32	0 / 0	0 / 0	0 / 0
64	2 / 0	1 / 0	1 / 0
128	9 / 0	4 / 5	4 / 4

Feasible / infeasible: 32: 0/32 · 64: 2/30 · 128: 16/16

Conditional challenge · 16 active + 8 shared controls

Mean absolute error · LUT4 ↓

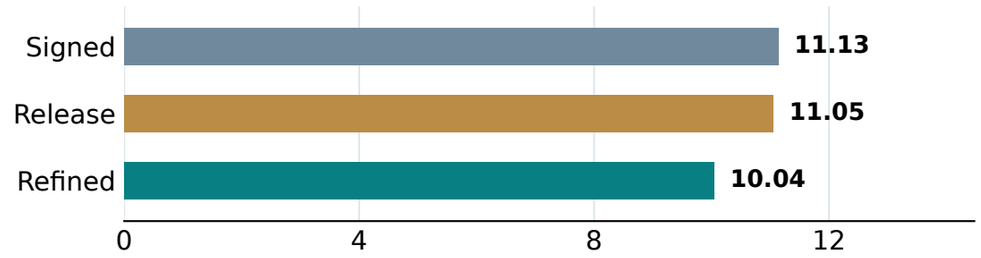

Positive-target MAPE · % ↓

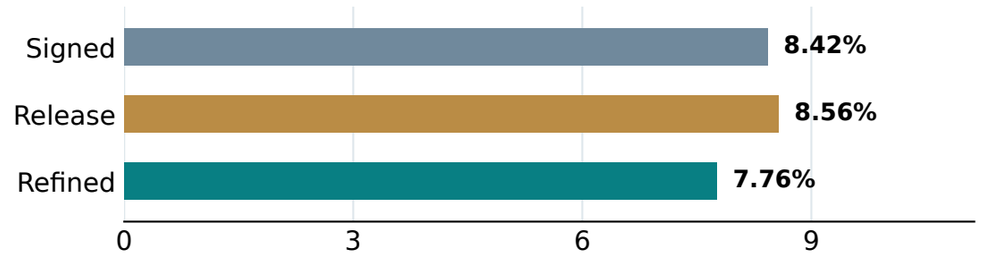

False accepts / false rejects · separate gates

Cap · LUT4	Signed	Release	Refined
32	0 / 0	0 / 0	0 / 0
64	1 / 0	0 / 0	0 / 0
128	3 / 0	2 / 16	2 / 0

Feasible / infeasible: 32: 0/24 · 64: 0/24 · 128: 21/3

Eight controls are shared between panels; no pooled metric or independent-cohort claim. Active-only gates are also audited.

The conditional sample is selected by frozen prediction disagreement, never labels. All pool features and predictions remain available.

Isolated iCE40 LUT4 estimation only; finite synthetic validation is not calibrated safety or a CPU speed, timing, power or application-transfer result.