What improved — and where it did not

Four separate matched comparisons. Lower MAE is better; cohorts, weighting and panel scales differ.

01

Mapping-aware descriptors

192 newly measured test circuits

Equal-stratum MAE · LUT4

Same 1,665 fitting circuits; 23 descriptors added.

02

Targeted training coverage

Previously consumed application panel

Equal-family MAE · LUT4

48 targeted circuits added; this is a development result.

03

Reserved-program transfer

327 originals · 521 program memberships

Equal-family / eligible-program MAE · LUT4

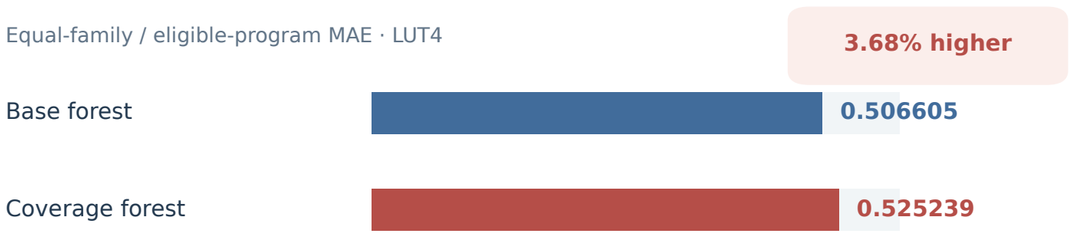

Five eligible families; empty recursion error remains undefined.

04

Scoped arithmetic stress

32 new related circuits · four templates

Equal-template MAE · LUT4

All 32 select the coverage forest; scoped policy equals it.

AVERAGE ERROR IS ONLY PART OF THE RESULT

Scoped stress: worst underprediction grows 2.83 → 32.01 LUT4.

Experimental policy; no conservative budget guarantee.

Saved comparisons: mapping descriptors · consumed applications · reserved source programs · covered-domain synthetic stress