V8 measurement-budget selection: instruction-count regret

Cap 64 LUT4 · measured feasibility required · lower regret is better · baseline saved instructions = 0

A Primary endpoint · one attempted candidate

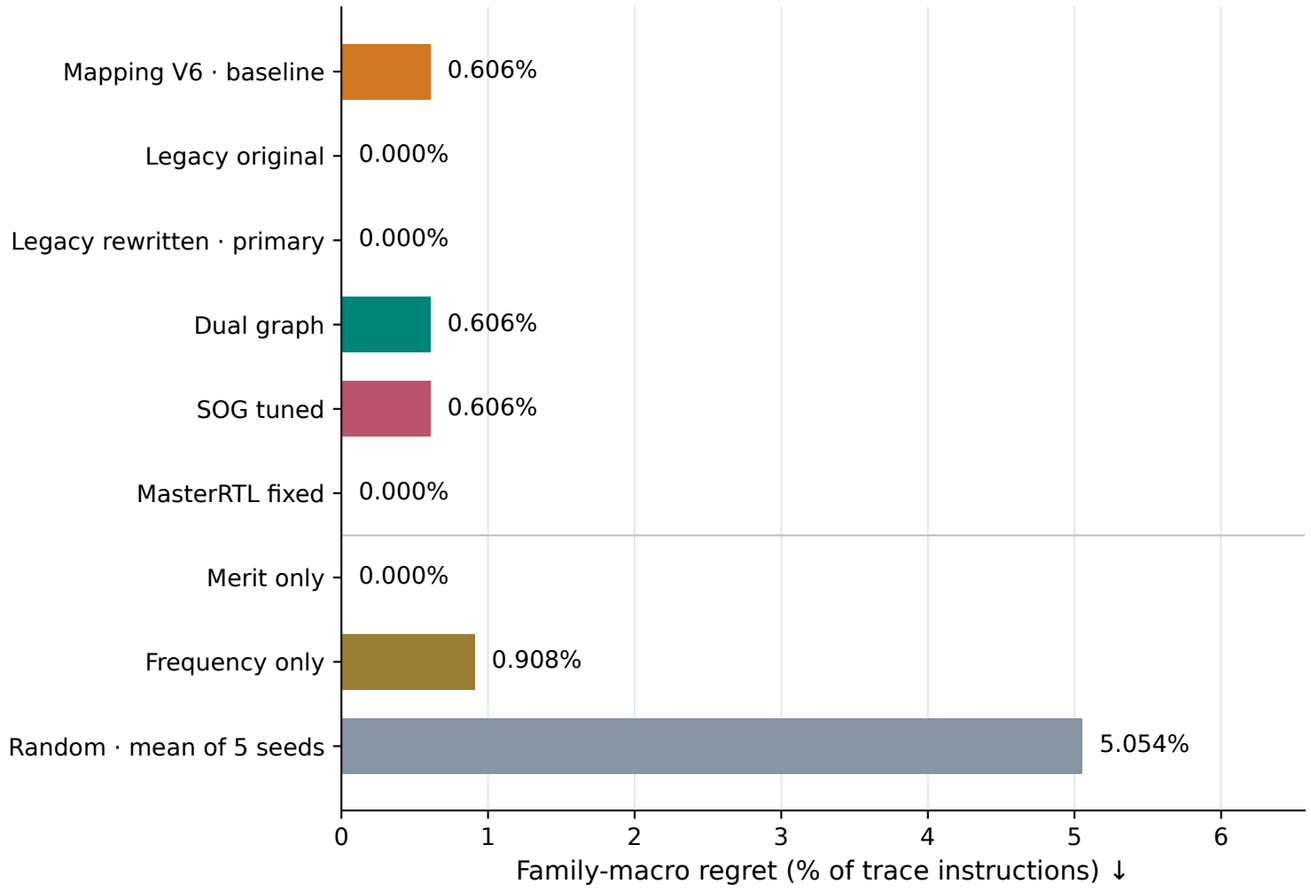

B Secondary budget sensitivity · same row order

1 attempt	3 attempts	5 attempts
0.606%	0.606%	0.606%
0.000%	0.000%	0.000%
0.000%	0.000%	0.000%
0.606%	0.606%	0.606%
0.606%	0.606%	0.606%
0.000%	0.000%	0.000%
0.000%	0.000%	0.000%
0.908%	0.908%	0.908%
5.054%	4.272%	3.758%

All values are regret as % of trace instructions.
Random averages all five fixed seeds; seeds are not population replicates.

Four primary source families; descriptive fixed-cohort results, no population confidence intervals. Benefit and regret use an overlap-safe retired-instruction-count proxy; no CPU speedup claim.
Six model policies prioritize predicted-feasible candidates by merit. Merit/frequency/random use the same candidate pools and attempt budgets.