V9: fixed training-coverage augmentation on consumed data

48 paid training labels + original 1,665 circuits → one fit on 1,713 rows · unchanged mapping217 features · ExtraTrees300

Consumed V8 application pools and V6 test regression · complete four-family comparison · no independent transfer claim

A Primary · complete V8 family MAE

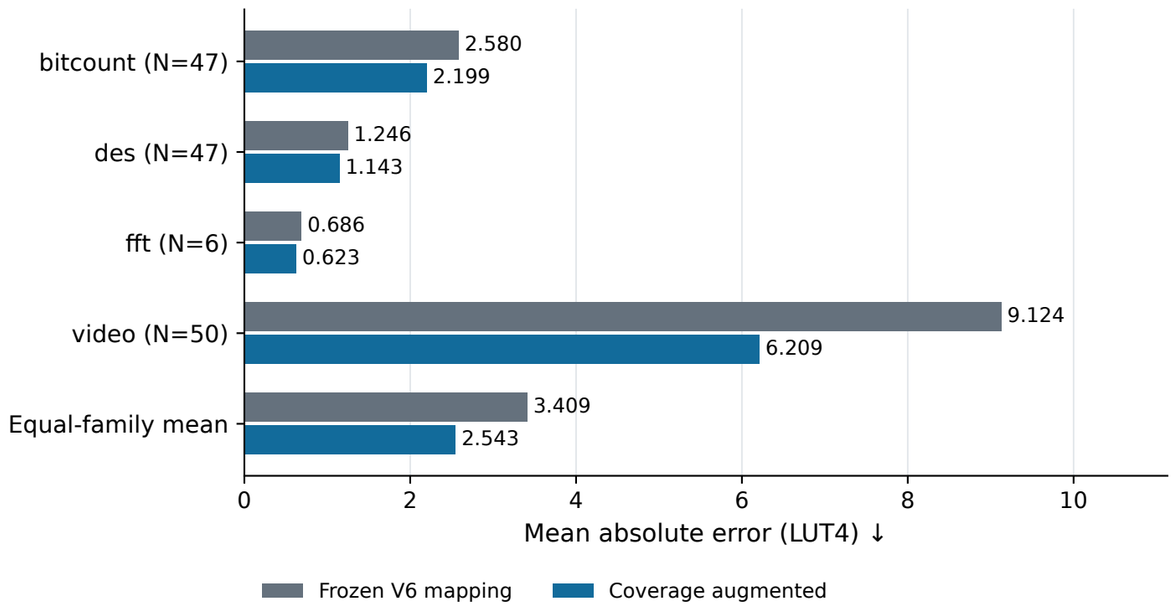

B Primary · paired family differences

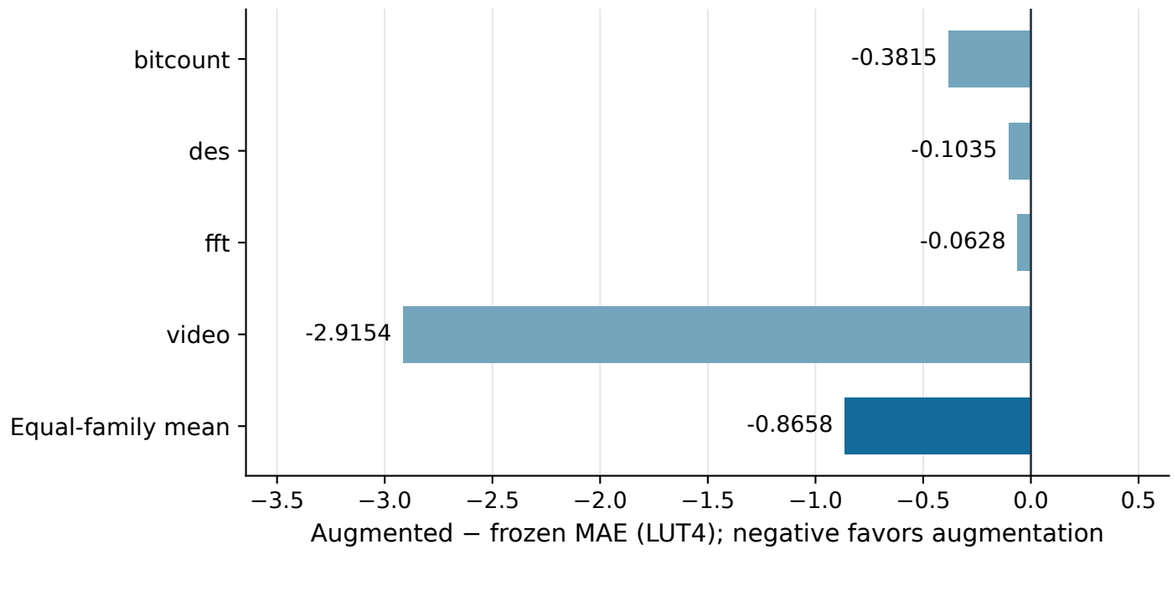

C Secondary · all six diagnosed constant multiplications

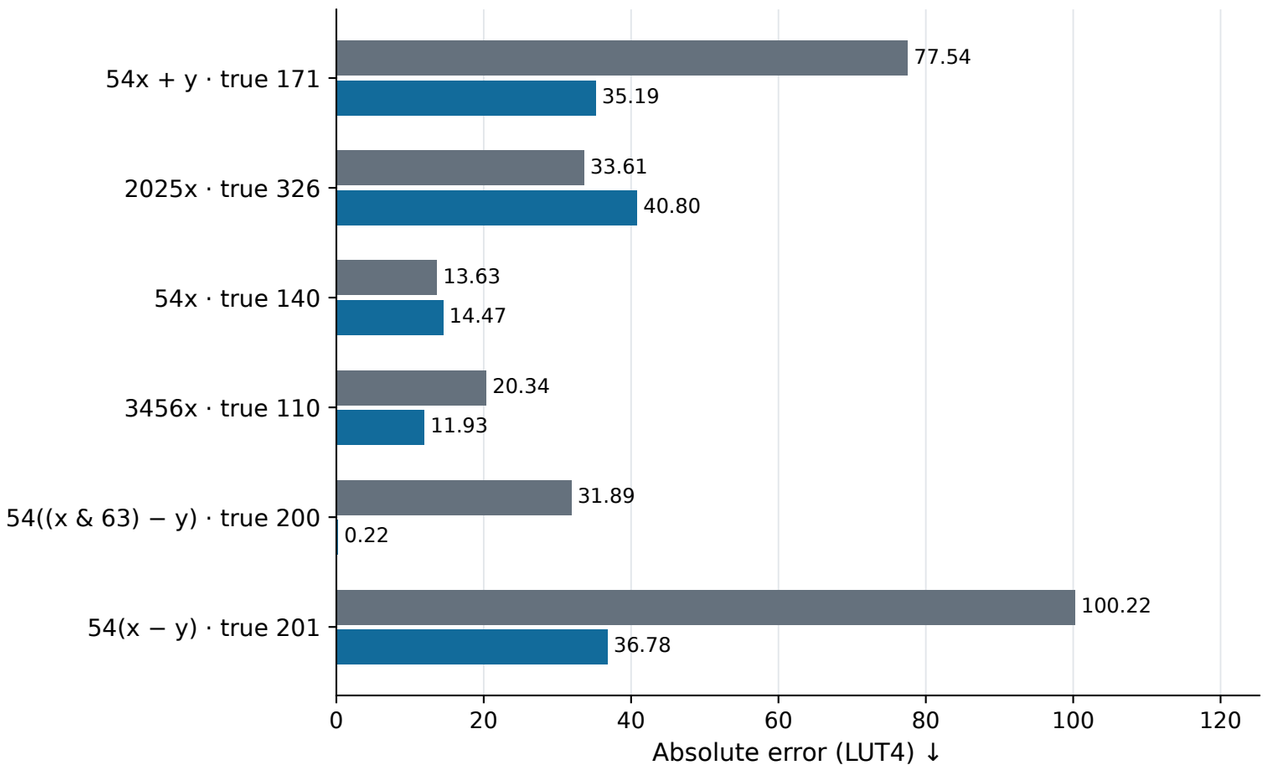

Colors match panel A. x=in0, y=in1; wrapping u32 arithmetic.
Full original formulas, hashes and points are in figure-data.json; all six cases shown.

D Secondary · complete consumed V6 test regression

Metric · all 192 cases	Frozen	Augmented
MAE · LUT4	32.1609	29.0538
Signed bias · LUT4	19.5400	17.0627
RMSLE	0.4688	0.4548
Positive MAPE · %	43.6295	41.3865
P95 absolute error · LUT4	145.6615	108.6384
P99 absolute error · LUT4	259.8994	264.3972
Worst absolute error · LUT4	1181.8467	1178.1167
Maximum underprediction · LUT4	94.8300	101.9600
Zero-target MAE · LUT4	3.7861	3.7628
Cap 32 · false accept / reject	4 / 14	4 / 14
Cap 64 · false accept / reject	6 / 5	6 / 5
Cap 128 · false accept / reject	1 / 4	1 / 6
Count · positive / zero	186 / 6	186 / 6
Under / over / exact	77 / 105 / 10	76 / 106 / 10

One fixed seed and one fit; all 48 coverage targets retained. Four-family descriptive development, no population confidence interval.
LUT4 prediction error is not ASIC physical area, instruction-selection quality, CPU speedup, timing or energy.