V12 · threshold and tail tradeoffs remain separate endpoints

All 1,713 OOF rows per view and seed · same five folds and zero override · raw/log contrasts are not an uncertainty or transfer claim

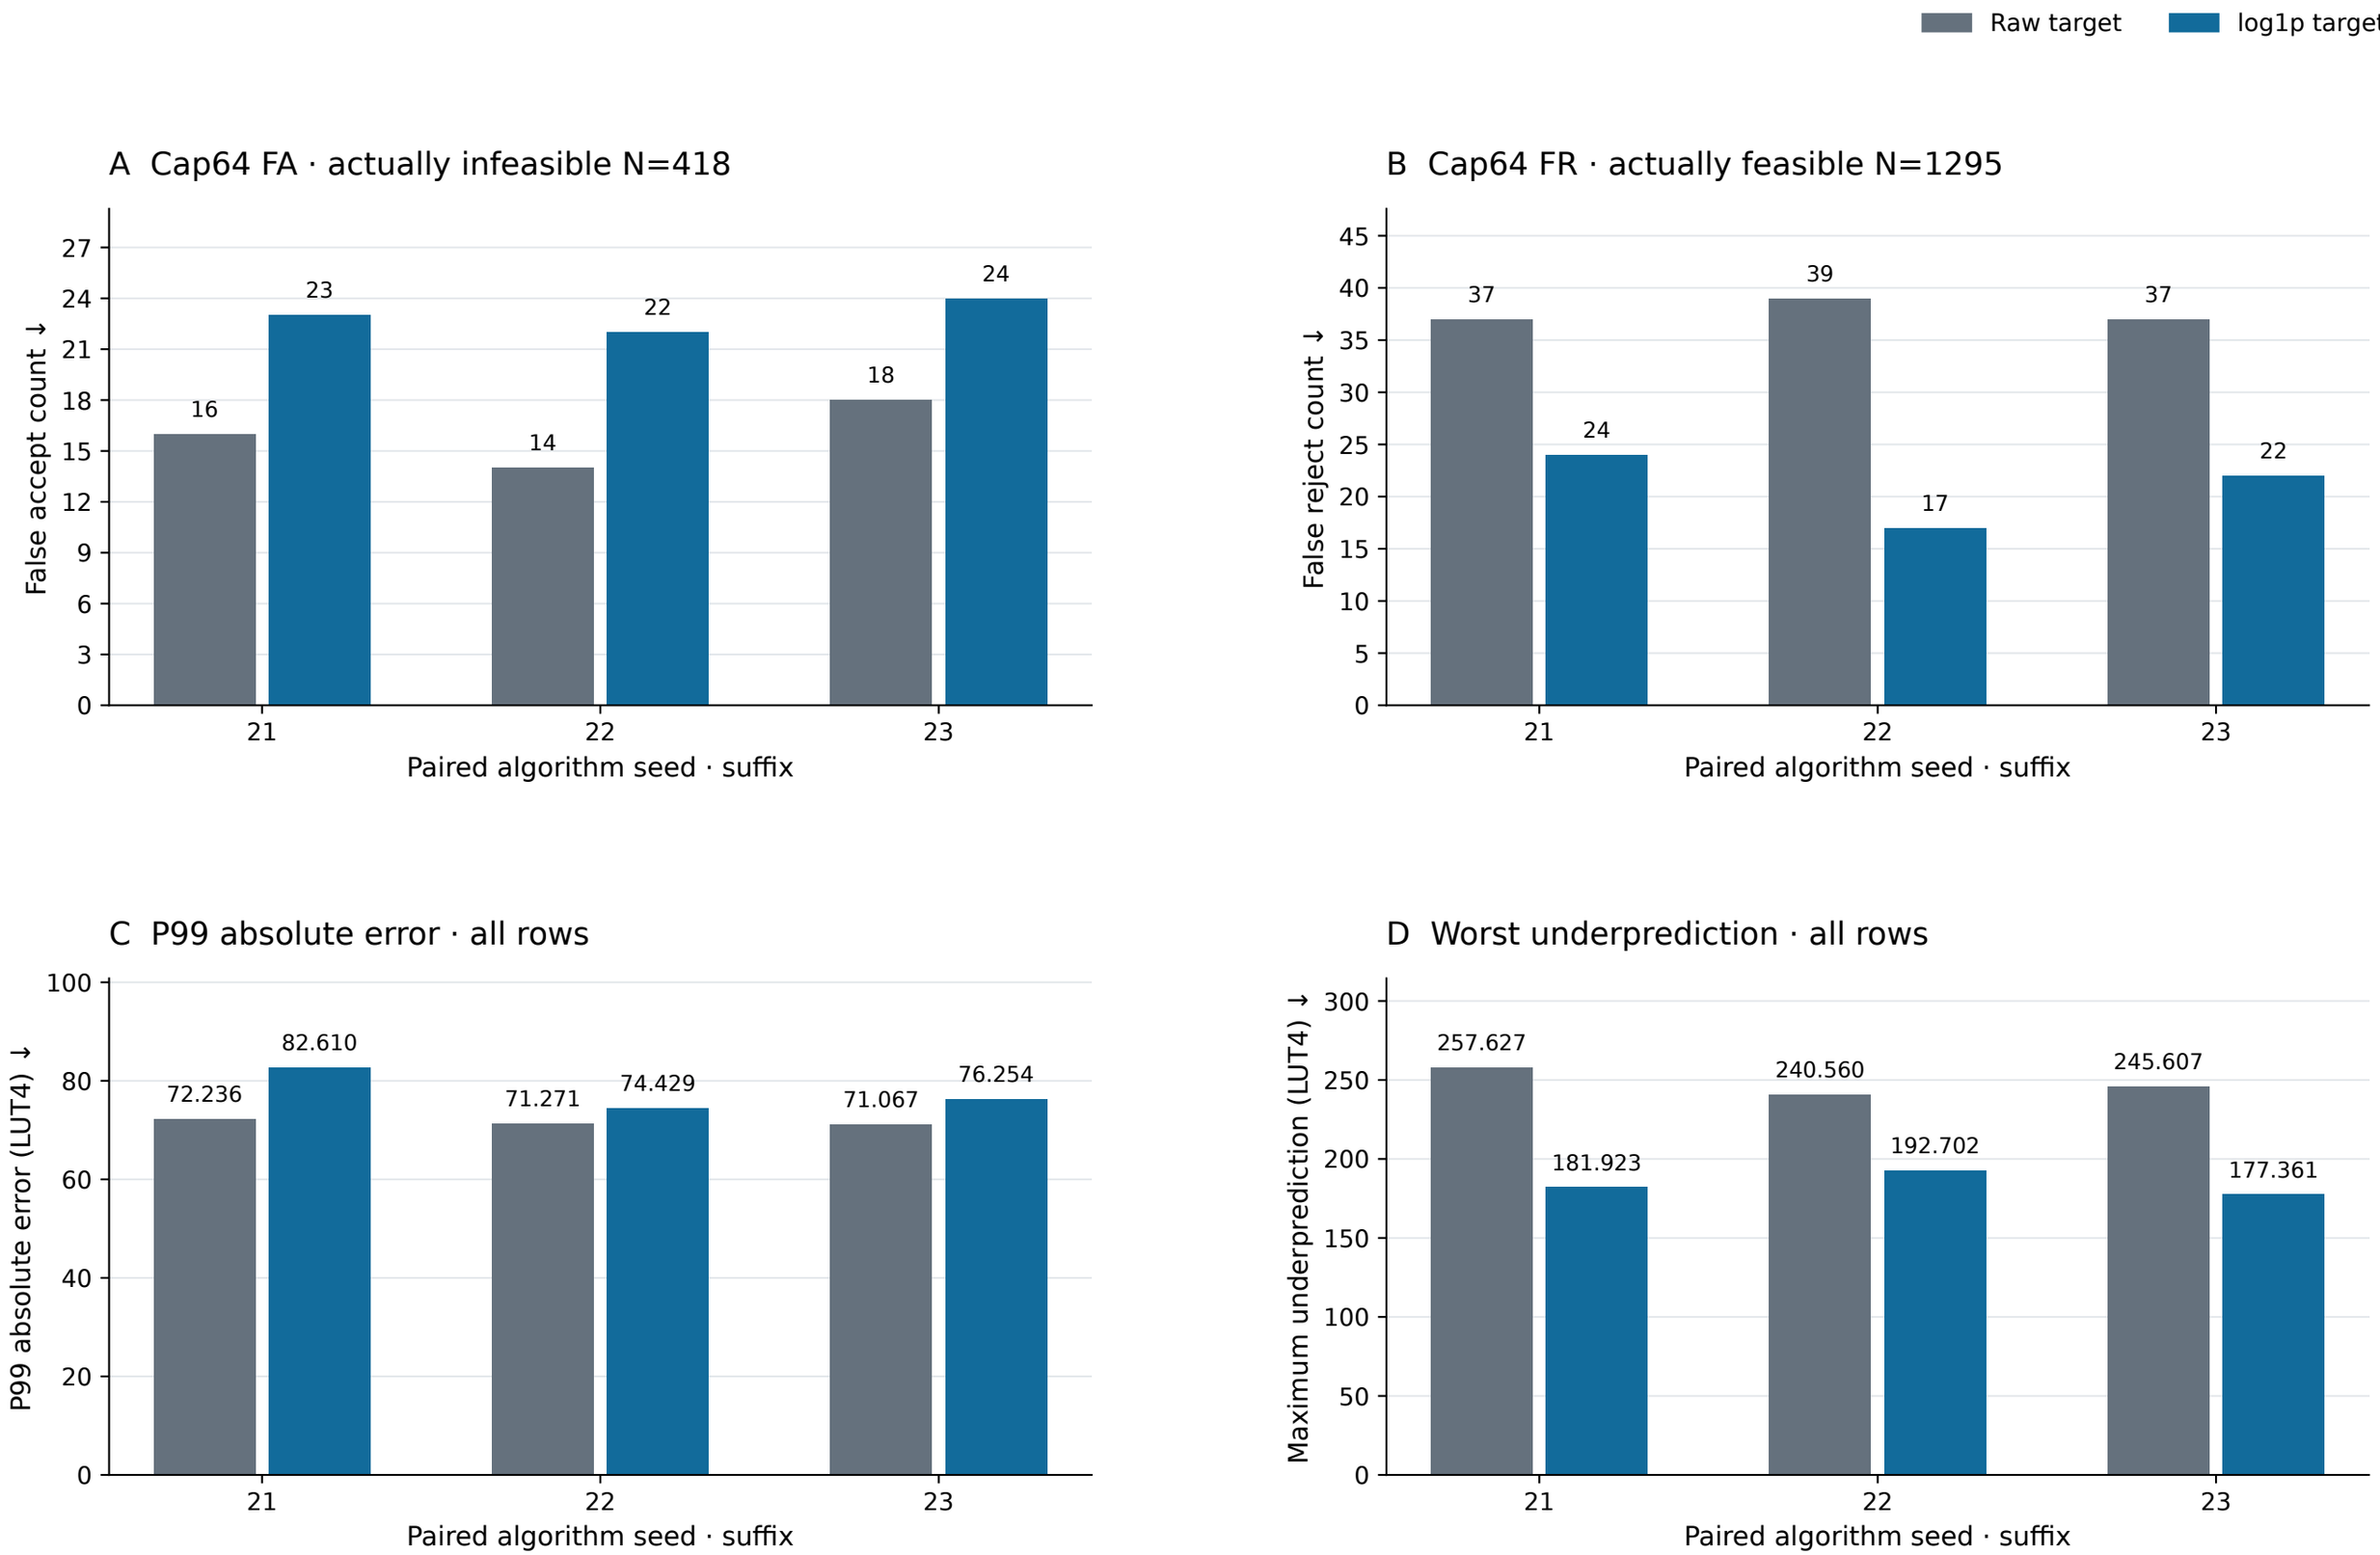

FA: prediction ≤64 when target >64. FR: prediction >64 when target ≤64. Equal-to-cap decisions count as feasible.
The log target increases cap64 false accepts and P99 error in every seed, while reducing false rejects and maximum underprediction.

Every paired seed is shown; no mean-error or one-sided-risk improvement substitutes for the other endpoints. Cap32/128 sensitivities and full score tables are retained in figure-data.json.
Values are static original-RTL LUT4 prediction errors on consumed fitting data, not hardware timing, packed area, board/CPU performance or population evidence.