V13 · budget errors and tails remain separate endpoints

All old 1,713 primary OOF rows in each view and seed · identical joint folds · cap64 boundary is inclusive

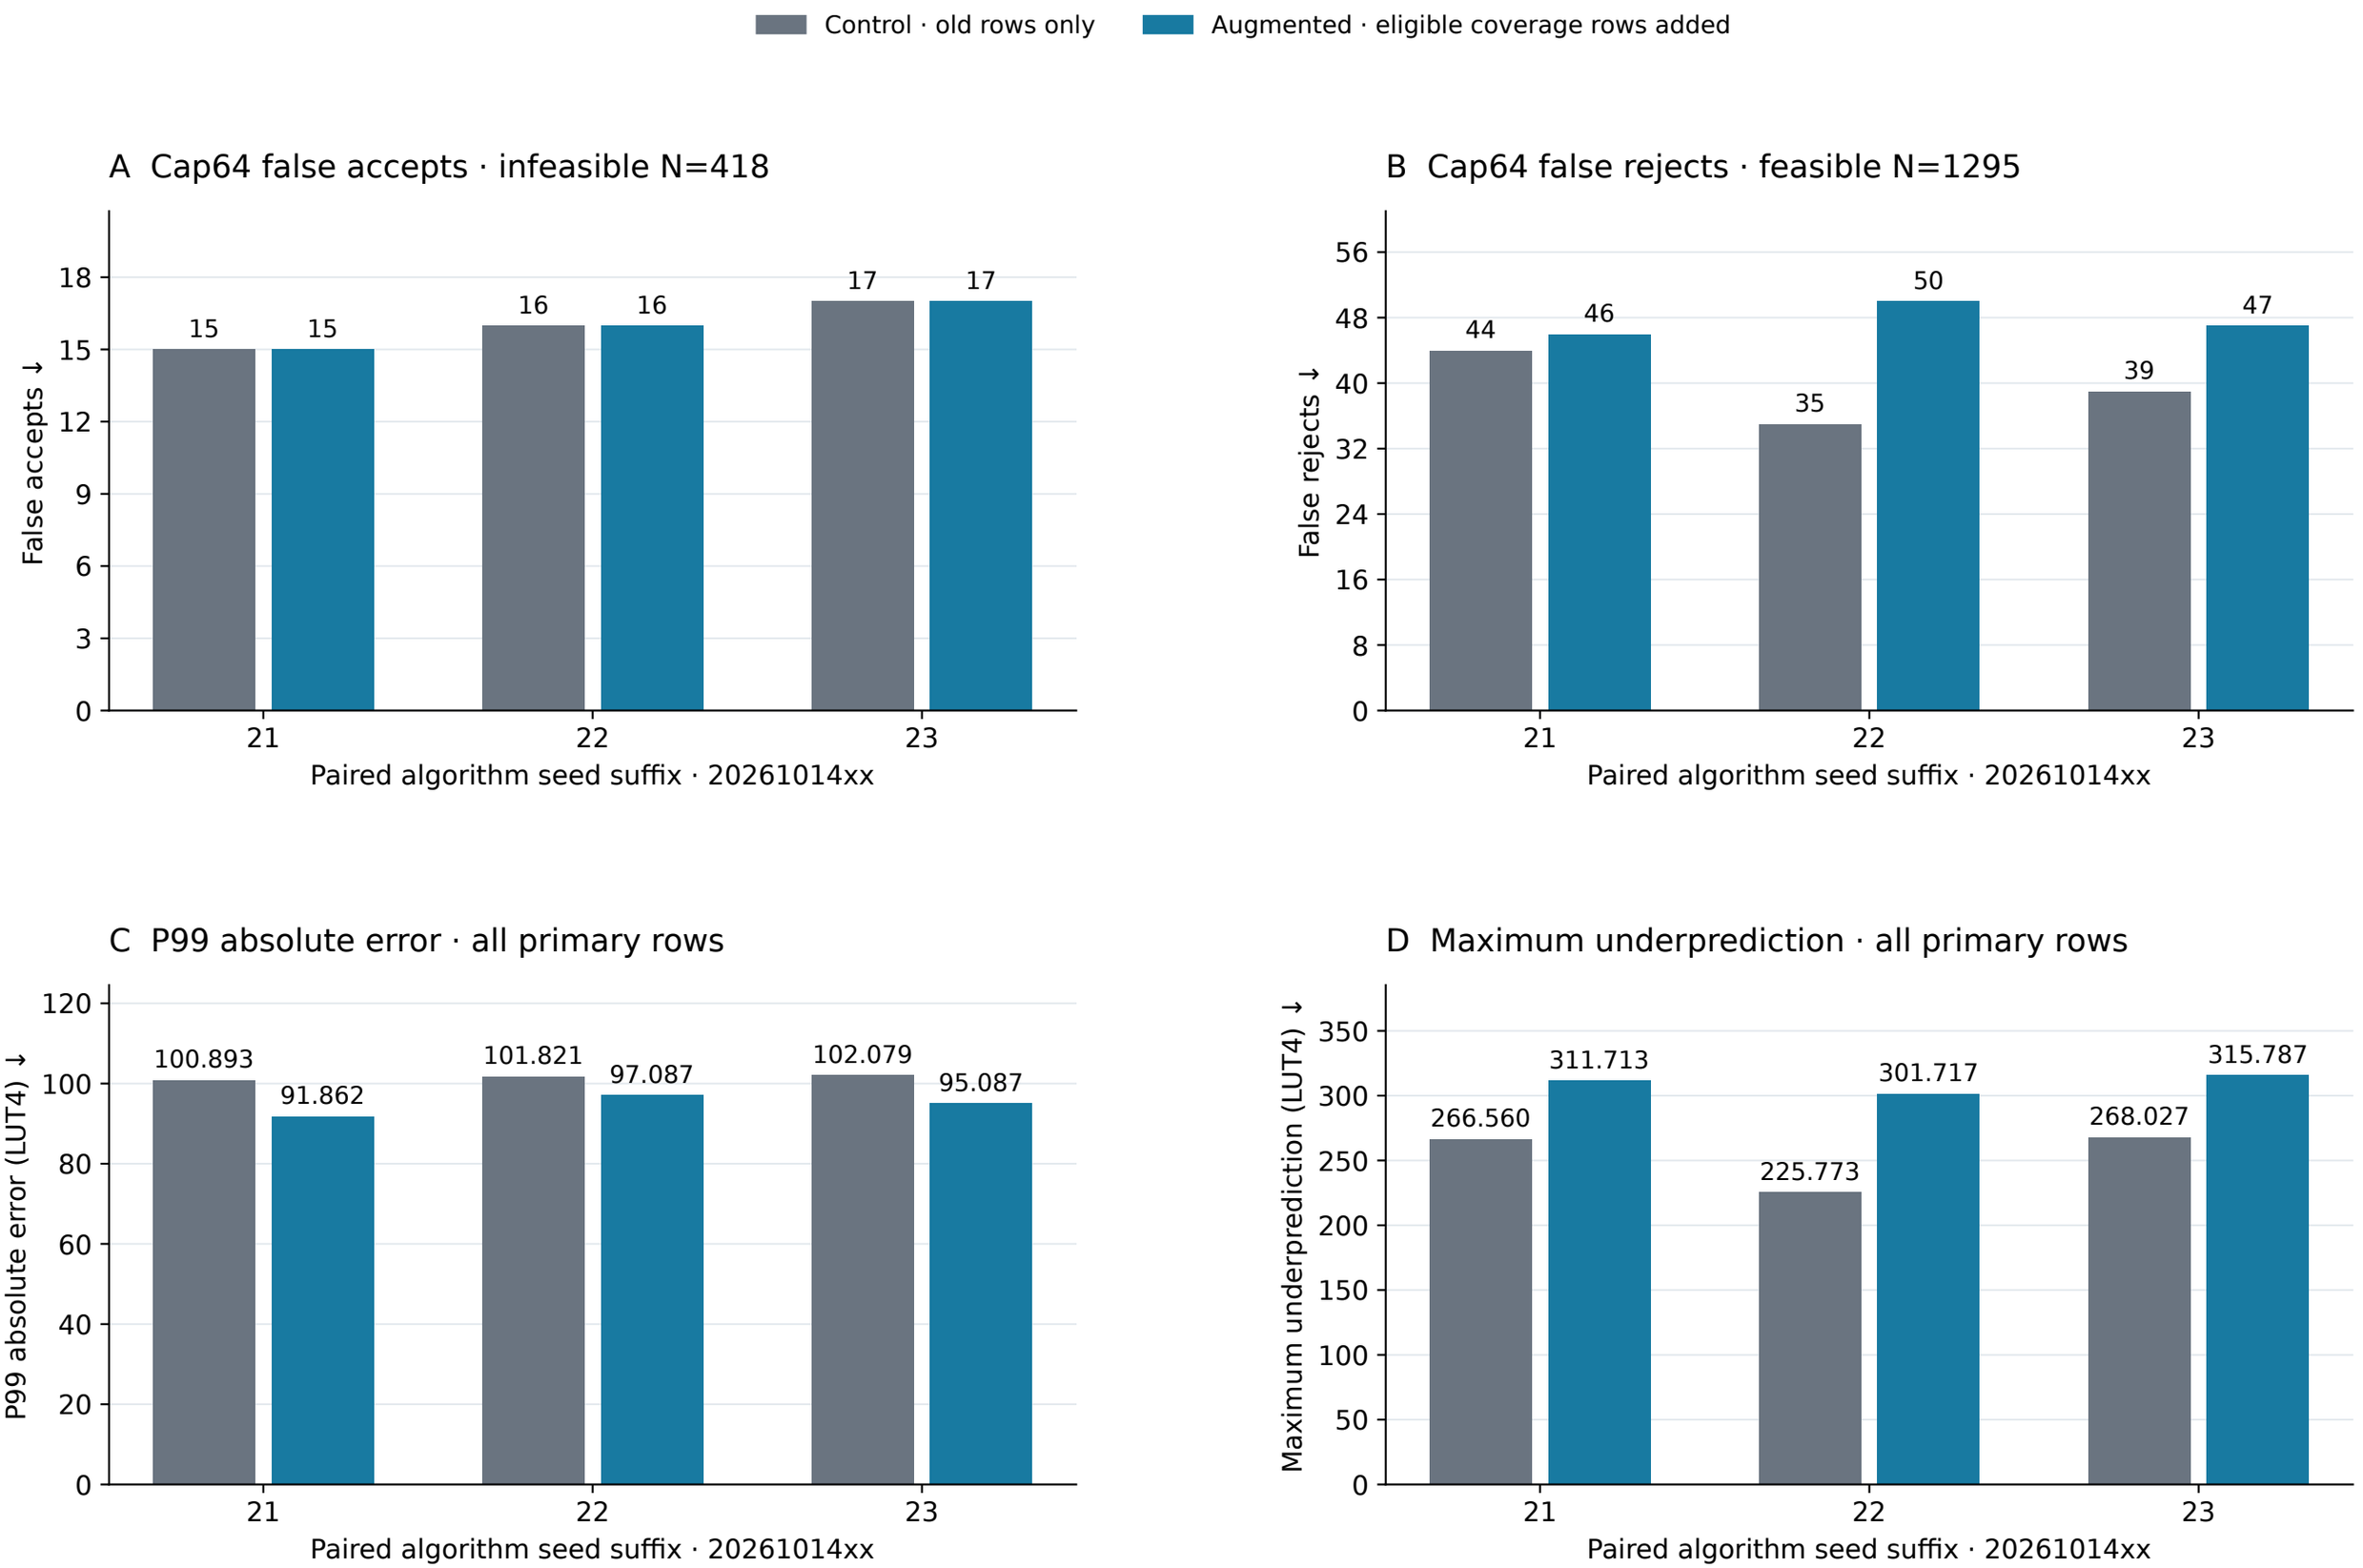

False accept: point ≤64 with target >64. False reject: point >64 with target ≤64. Lower is better for each endpoint separately.
Three paired algorithm seeds are displayed in full; an improved mean does not erase a worse cap decision or tail. Cap32/128 sensitivities remain in figure-data.json.

Static original-RTL LUT4 prediction errors on consumed development data. No statistical confidence intervals, calibrated safety bound, hardware timing, power or CPU-performance inference.
Both views retain the same original-generic-graph zero override; canonical/accepted aliases, exact selected-vector collisions and complete contexts close the shared fold map.